

The graph displays the training loss for three different models over 5000 epochs. The 'Baseline loss curve' (gray) and 'Model 0 loss curve' (blue) are highly volatile, fluctuating between approximately 0.01 and 0.09. The 'Ensemble data loss' (blue line with markers) is a smoothed version of the ensemble's performance, starting at approximately 0.035 at epoch 500 and gradually increasing to about 0.065 by epoch 5000.

Epoch	Baseline loss curve (approx.)	Model 0 loss curve (approx.)	Ensemble data loss (approx.)
0	0.05	0.05	-
500	0.03	0.03	0.035
1000	0.04	0.04	0.055
1500	0.03	0.03	0.058
2000	0.04	0.04	0.060
2500	0.03	0.03	0.062
3000	0.04	0.04	0.063
3500	0.03	0.03	0.064
4000	0.04	0.04	0.065
4500	0.03	0.03	0.066
5000	0.04	0.04	0.065