

The graph displays three loss curves over 4000 epochs. The y-axis represents the loss, ranging from 0.00 to 0.10. The x-axis represents the epoch number, ranging from 0 to 4000.

- Baseline loss curve (grey line):** This curve starts at a high loss of approximately 0.10 and quickly drops to a noisy plateau around 0.05 by epoch 500. It remains relatively stable with high-frequency noise throughout the training process.
- Model 0 loss curve (solid blue line):** This curve starts at a high loss of approximately 0.10 and drops sharply to about 0.05 by epoch 500. It then continues to decrease more gradually, reaching a loss of approximately 0.01 by epoch 4000.
- Ensemble data loss (blue line with circular markers):** This curve follows the Model 0 loss curve closely, starting at approximately 0.10 and dropping to about 0.05 by epoch 500. It then continues to decrease, reaching a loss of approximately 0.01 by epoch 4000. The markers indicate specific data points or checkpoints.

The graph illustrates that the ensemble data loss is significantly lower than the baseline loss curve, suggesting improved performance when using the ensemble data.