

The graph illustrates the performance of three different loss functions over 4000 epochs. The y-axis represents the loss, ranging from 0.00 to 0.10. The x-axis represents the epoch number, ranging from 0 to 4000.

- Baseline loss curve (grey line):** This curve starts at a high loss of approximately 0.10 and quickly drops to a noisy plateau around 0.05 by epoch 400. It remains relatively stable with high-frequency noise throughout the training process.
- Model 0 loss curve (blue line):** This curve starts at a high loss of approximately 0.10 and decreases steadily to about 0.05 by epoch 400. After epoch 400, it remains relatively flat, ending at approximately 0.05 at epoch 4000.
- Ensemble data loss (blue line with dots):** This curve starts at a high loss of approximately 0.10 and decreases steadily and more rapidly than the other two curves, reaching a final loss of approximately 0.01 by epoch 4000.

The Ensemble data loss curve shows the most significant improvement in performance, achieving the lowest loss value by the end of the training process.